

# INTRODUCTION

## CS3021/3421 Computer Architecture II

Dr Jeremy Jones

Office: O'Reilly Institute F.11

Email: [jones@scss.tcd.ie](mailto:jones@scss.tcd.ie)



## INTRODUCTION

BACS/MCS and BAI/MAI

students

# INTRODUCTION

## BACS/MCS YEAR 3 TIMETABLE

School of Computer Science and Statistics  
Integrated Computer Science: Year 3 Timetable 2016-17

Time	Monday	Tuesday	Wednesday	Thursday	Friday
09.00 – 10.00			MT: CS3012: Lect LB04		
10.00 – 11.00	MT: CS3021: Lect LB08 HT: CS3031: Lect LB08	MT: CS3071: Lect LB01 HT: CS3031: Lect LB01	MT: CS3011: Lect LB01 HT: CS3061: Lect LB04	MT: CS3021: Lect LB08 HT: CS3031: Lect LB08	HT: CS3081: Lect LB01
11.00 – 12.00					
12.00 – 13.00	MT: CS3012: Lect Joly	HT: CS3061: Lect LB04		MT: CS3071: Lect LB01	
13.00 – 14.00	MT: CS3011: Lect LTEE1	MT: CS3012: Lect Goldhall HT: CS3013: Lect LB01 (2hrs)	HT: ST3009: Lab: LG12/35/36		
14.00 – 15.00	MT: CS3016 Lect LB01/ICT1/2 HT: CS3061 Lect LB04	MT: CS3011: Lect LG12 HT: CS3013: Lect LB01	MT: CS3041: Lect LB01	MT: CS3016: Lect LB04/ICT1 HT: CS3014: Lect LB04/ICTLab2 (2hrs)	HT: CS3081: Tut: LB01
15.00 – 16.00			MT: CS3021: Lect M17 HT: ST3009: Lect: LB01 (2hrs)	HT: CS3014: Lect LB04	MT: CS3016: Lect LB01
16.00 – 17.00	MT: CS3011: Lab LG12	HT: CS3014: Lect LB04	HT: ST3009: Lect: LB01 MT: CS3071: Lab ICT Lab1/2	MT: CS3041: Lect LB08	
17.00 – 18.00					

# INTRODUCTION

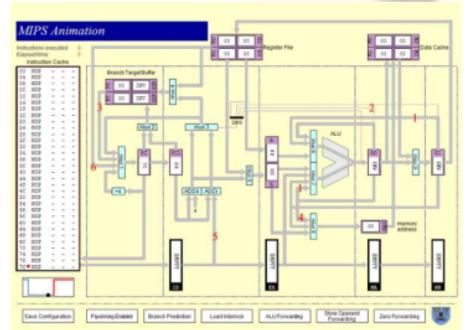
## BAI/MAI YEAR 4 TIMETABLE

SENIOR SOPHISTER ENGINEERING, 2016/17 - ELECTRONIC AND COMPUTER ENGINEERING										
Date of issue: 22nd September, 2016										
DAY	SEMESTER	0900 - 1000	1000 - 1100	1100 - 1200	1200 - 1300	1300 - 1400	1400 - 1500	1500 - 1600	1600 - 1700	1700 - 1800
MONDAY	First semester	4C4 [M20]	CS3421 [LB08]		4C5 [M20]		4C4 [M20]	CS4052 [LB04]	4C7 [M20]	4E1 [HAM4]
	Second semester	4B9 [M17]	4C2 [M17]	CS7434 [LB01]	4C15 [M21]		Electronic/Computer Engineering projects			CS7453 [LB04]
TUESDAY	First semester	4C5 [M17]	CS4053 [LB08]	CS4052 [LB04/ICTLAB1/ICTLAB2]		CS4053 [LB04 and ICTLAB2]		4C8 [M21]	4E1 [HAM4]	
	Second semester		4C2 [M21]	CS4D2B [M21]	4C15 [M21]	4B9 tutorial [LB04]	4B9 laboratories [EE LABS]	4B9 [2037]		
WEDNESDAY	First semester	4C8 [EE LABS]			4C5 [M20]	CS4053 [LB01 and ICTLAB1]	CS4D2A [LB01]	CS3421 [M17]	4C8 [M21]	4C4 [M21]
	Second semester	CS4D2B [M20]		4B9 [2039]	4C15 [PBSR1]	CS7434 [LB01 and ICTLAB2]	CS4D2B [LB01]	CS7453 [HAM1]		
THURSDAY	First semester	4C7 [M20]	CS3421 [LB08]		CS4031 [LB04]		4C7 [M17]	4C8 [HAM5]	CS4D2A [LB08]	
	Second semester		4C2 [M17]		4C15 [M21]			4C2 [CLT]	CS4D2B [LB08]	
FRIDAY	First semester	Electronic Engineering laboratories/projects (C and CD stream) [EE LABS]					4C5 [M20]	4C7 [M21]	4C4 [M21]	
	Second semester	Electronic Engineering laboratories/projects (C and CD stream) [EE LABS]								

# INTRODUCTION

## SYLLABUS

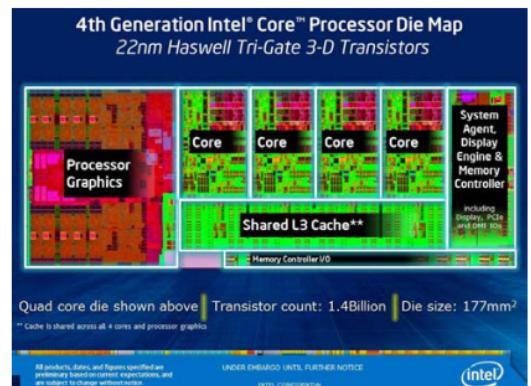
- IA32 and x64 assembly language programming
- IA32 and x64 procedure calling conventions
- RISC vs CISC
- RISC-1 design criteria and architecture
- Register windows and delayed jumps
- Instruction level pipelining
- DLX/MIPS pipeline
- Resolving data, load and control hazards
- Virtual Memory
- Memory management units [MMUs]
- Multi-level page tables and TLBs
- MMU integration with an OS



# INTRODUCTION

## SYLLABUS ...

- Cache organisation (L, K and N)
- Cache operation and performance
- The 3 Cs
- Virtual vs physical caches
- Pseudo LRU and LRU replacement policies
- Address trace analysis
- Multiprocessor architectures
- Cache coherency
- Cache coherency protocols [write-through, write-once, Firefly and MESI]



# INTRODUCTION

## ASSESSMENT

Coursework: **20%**

- 5 or 6 tutorials + a coursework project



Examination: **80%**

- answer 3 out of 4 questions in 2 hours

# INTRODUCTION

## COURSE WEB PAGE

<https://www.scss.tcd.ie/Jeremy.Jones/CS3021/CS3021.htm>

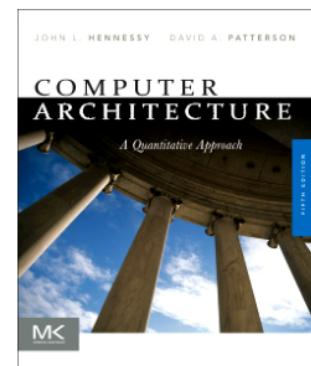
For

- Lecture notes
- Tutorials (questions, answers and marks)
- Coursework
- Miscellaneous materials

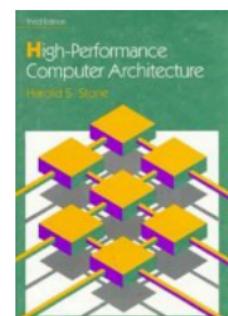
# INTRODUCTION

## Useful Books

Computer Architecture - a Quantitative Approach  
John Hennessy and David Patterson



High Performance Computer Architecture  
Harold S. Stone



# INTRODUCTION

Get Started on Wednesday @ 3pm M17

See you there!

# IA32 AND X64

## Microprocessor Design Trends

- Joy's Law [Bill Joy of BSD4.x and Sun fame]

$$\text{MIPS} = 2^{\text{year}-1984}$$

- Millions of instructions per second [MIPS] executed by a single chip microprocessor
- More realistic rate is a doubling of MIPS every 18 months [or a quadrupling every 3 years]
- What ideas and techniques in new microprocessor designs have contributed to this continued rate of improvement?

## IA32 AND X64

Some of the ideas and techniques used...

- smaller VLSI feature sizes [1 micron ( $\mu$ ) -> 10nm]
- increased clock rate [1MHz -> 4GHz] 
- reduced vs complex instruction sets [RISC vs CISC]
- burst memory accesses
- integrated on-chip MMUs, FPUs, ... *used to be on separate chips*
- pipelining
- superscalar [multiple instructions/clock cycle]
- multi-level on-chip instruction and data caches
- streaming SIMD [single instruction multiple data] instruction extensions [MMX, SSEx]
- multiprocessor support
- hyper threading and multi core
- direct programming of graphics co-processor
- high speed point to point interconnect [Intel QuickPath, AMD HyperTransport]
- solid state disks

- Speed of main memory access still the bottleneck

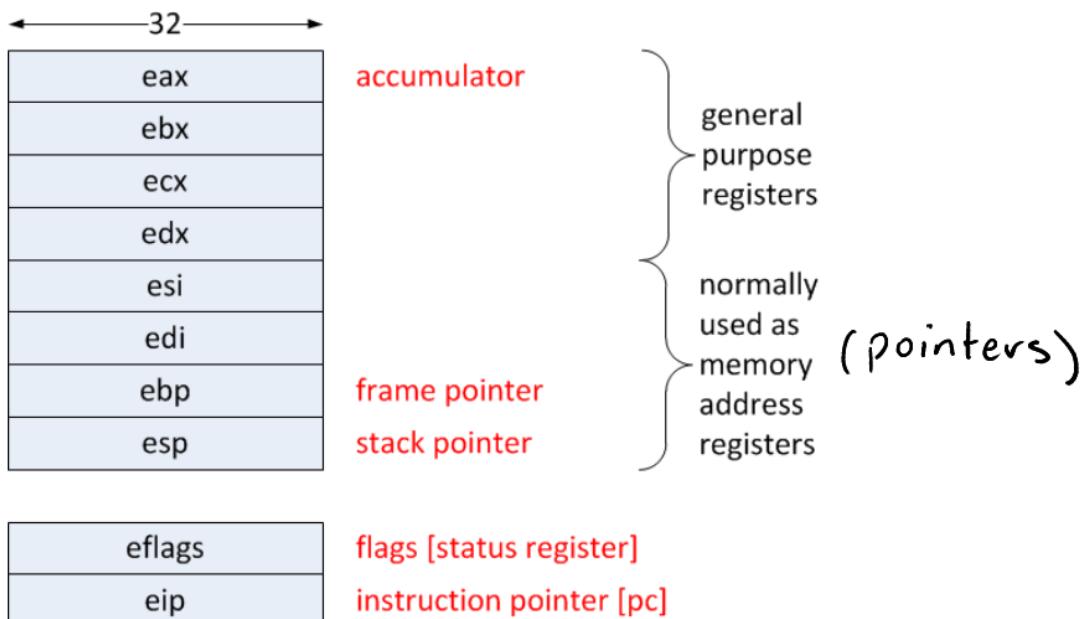
# IA32 AND X64

## IA32 [Intel Architecture 32 bit]

- IA32 first released in 1985 with the 80386 microprocessor
- IA32 still used today by current Intel CPUs
- modern Intel CPUs have many additions to the original IA32 including MMX, SSE1, SSE2, SSE3, SSE4 and SSE5 [**Streaming SIMD Extensions**] and even an extended 64 bit instruction set when operating in 64 bit mode [**named IA-32e or IA-32e or x64**]
- 32 bit CPU [**performs 8, 16 and 32 bit arithmetic**]
- 32 bit virtual and physical address space  $2^{32}$  bytes [**4GB**]
- each instruction a multiple of bytes in length [**1 to 17+**]

# IA32 AND X64

Registers [not as many as a typical RISC]

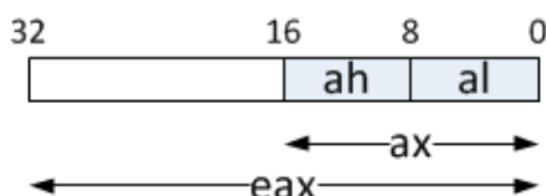


NB: floating point and SSE registers etc. not shown

# IA32 AND X64

## Registers...

- "e" in eax = extended = 32bits



- possible to access 8 and 16 bit parts of eax, ebx, ecx and edx using alternate register names

# IA32 AND X64

## Instruction Format

- two address [will use Microsoft assembly language syntax used by VC++, MASM]

add  $\xrightarrow{\text{Source}}$  eax, ebx ; eax = eax + ebx [right to left]  
*Source and destination*

- alternative gnu syntax

$\xrightarrow{\text{Left}}$  addl %ebx, %eax ; eax = eax + ebx [left to right]

- two operands normally

$\xrightarrow{\text{register/register}}$   
register/immediate  
register/memory  
memory/register

- memory/memory and memory/immediate **NOT** allowed

# IA32 AND X64

## Supported Addressing Modes

[*a*] = contents of  
memory address *a*

<i>addressing mode</i>	<i>example</i>	
immediate	mov eax, <i>n</i>	eax = <i>n</i>
register	mov eax, ebx	eax = ebx
direct/absolute	mov eax, [ <i>a</i> ]	eax = [ <i>a</i> ]
indexed	mov eax, [ebx]	eax = [ebx]
indexed	mov eax, [ebx+ <i>n</i> ]	eax = [ebx + <i>n</i> ]
scaled indexed	mov eax, [ebx* <i>s</i> + <i>n</i> ]	eax = [ebx* <i>s</i> + <i>n</i> ]
scaled indexed	mov eax, [ebx+ecx]	eax = [ebx + ecx]
scaled indexed	mov eax, [ebx+ecx*s+n]	eax = [ebx + ecx*s + <i>n</i> ]

- address computed as the sum of a register, a scaled register and a 1, 2 or 4 byte signed constant *n*; can use most registers
- scaled indexed addressing used to index into arrays
- scaling constant *s* can be 1, 2, 4 or 8;

# IA32 AND X64

## Assembly Language Tips

- size of operation can often be determined implicitly by assembler, but when unable to do so, size needs to be specified explicitly

mov eax, [ebp+8] ; implicitly 32 bit [as eax is 32 bits]

mov ah, [ebp+8] ; implicitly 8 bit [as ah is 8 bits]

dec [ebp+8] ; decrement memory location [ebp+8] by 1  
; assembler unable to determine operand size  
; is it an 8, 16 or 32 bit value??

= this is a pointer to a 32 bit value

dec DWORD PTR [ebp+8] ↵ ; make explicitly 32 bit

dec WORD PTR [ebp+8] ; make explicitly 16 bit

dec BYTE PTR [ebp+8] ; make explicitly 8 bit

NB: unusual assembly language syntax

# IA32 AND X64

## Assembly Language Tips...

- memory/immediate operations NOT allowed

~~mov [ebp+8], 123~~ ; NOT allowed and operation size ALSO unknown

mov eax, 123 ; use 2 instructions instead...  
mov [ebp+8], eax ; explicitly 32 bits

- lea [**load effective address**] is useful for performing simple arithmetic

\* lea eax, [ebx+ecx\*4+16] ; eax = ebx+ecx\*4+16

# IA32 AND X64

## Basic Instruction Set

		push	push onto stack
		pop	pop from stack
mov	move		
xchg	exchange Swap Reg values	sar	shift arithmetic right
add	add	shl	shift logical left
sub	subtract	shr	shift logical right
imul	signed multiply		
mul	unsigned multiply	jmp	unconditional jump
inc	increment by 1	j {e, ne, l, le, g, ge}	signed jump
dec	decrement by 1	j {b, be, a, ae}	unsigned jump
neg	negate		
cmp	compare	call	call subroutine
lea	load effective address	ret	return from subroutine
test	AND operands and set flags		

and	and
or	or
xor	exclusive or
not	not

- should be enough instructions to complete tutorials
- Google [Intel® 64 and IA-32 Architectures Software Developer's Manual 2A, 2B, 2C](#) for details

use to test before branching

arithmetic shift divides by 2.

e = equal

b = below

a = above

 [banh](#) said:  
19th January 2006 14:50

**difference between arithmetic and logical shift**

logical shift: 0 is shifted in

arithmetic shift: the sign is preserved

# IA32 AND X64

## Assembly Language Tips...

- quickest way to clear a register?

```
xor    eax, eax           ; exclusive OR with itself
```

```
mov    eax, 0             ; instruction occupies more bytes and...
                           ; probably takes longer to execute
```

- quickest way to test if a register is zero?

```
test   eax, eax           ; AND eax with itself, set flags and...
je     ...                 ; jump if zero
```

# IA32 AND X64

## Function Calling

remainder of the steps normally carried out during a function/procedure call and return

- pass parameters [evaluate and push on stack]
  - enter new function [push return address and jump to first instruction of function]
  - allocate space for local variables [on stack by decrementing esp]
  - save registers [on stack]
- current PC - ebp*  
*branch to ...*

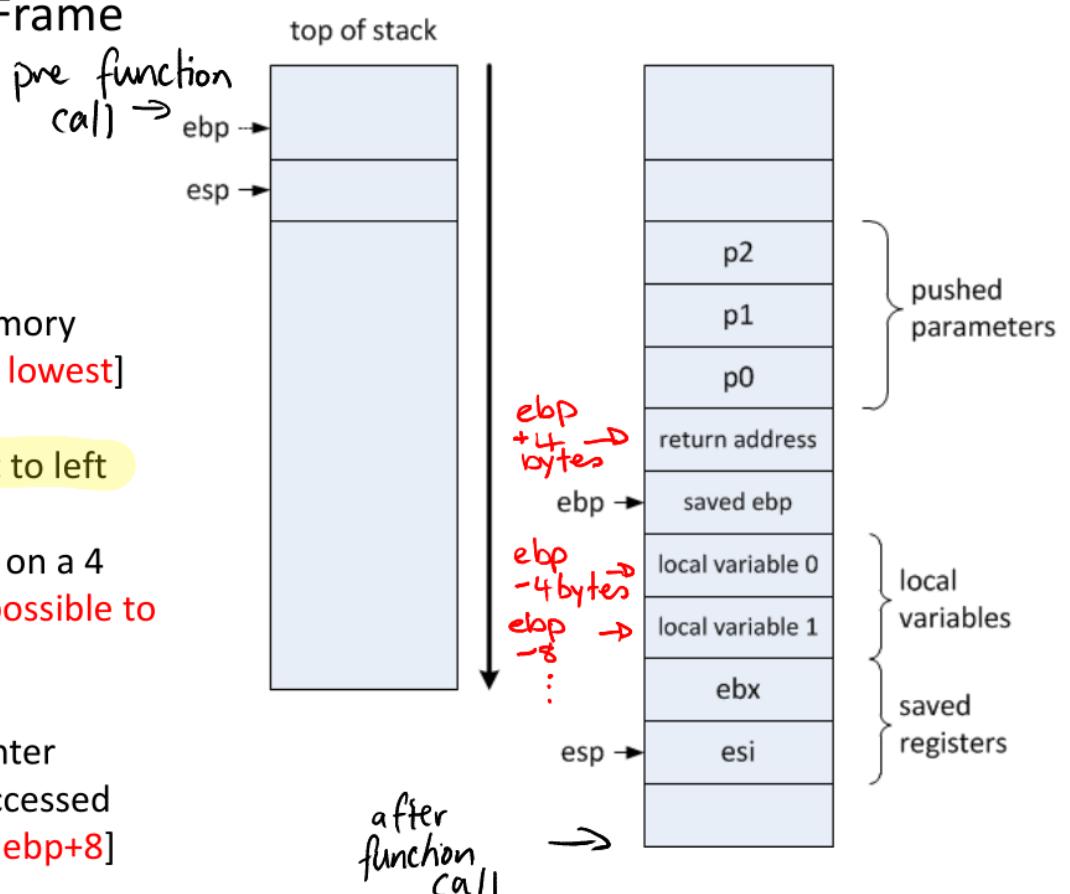
*<function body>*

- restore saved registers [from stack]
- de-allocate space for local variables [increment esp]
- return to calling function [pop return address from stack]
- remove parameters [increment esp]

# IA32 AND X64

## IA32 Function Stack Frame

- stack frame after call to  $f(p_0, p_1, p_2)$
- stack grows down in memory [from highest address to lowest]
- parameters pushed right to left
- NB: stack always aligned on a 4 byte boundary [it's not possible to push a single byte]
- ebp used as a frame pointer parameters and locals accessed relative to ebp [eg  $p_0 @ \text{ebp}+8$ ]



# IA32 AND X64

## IA32 Calling Conventions

- several IA32 procedure/function calling conventions
- use Microsoft \_cdecl calling convention [as per previous diagram] so C/C++ and IA32 assembly language code can mixed

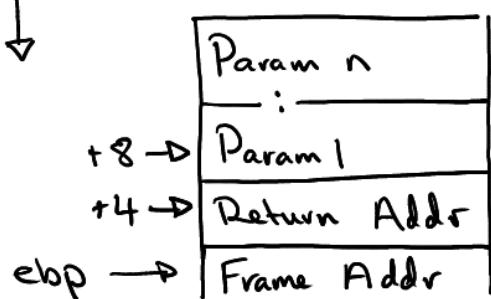
function result returned in eax *Not saved b/w function calls for this reason*

eax, ecx and edx considered volatile and are NOT preserved across function calls

*\* caller knows how many parameters  
caller removes parameters were passed*

- why are parameters pushed right-to-left??

C/C++ pushes parameters right-to-left so functions like `printf(char *formats, ...)` [which can accept an arbitrary numbers of parameters] can be handled more easily since the first parameter is always stored at [ebp+8] irrespective of how many parameters are pushed



# IA32 AND X64

## Accessing Parameters and Local Variables

- ebp used as a frame pointer; parameters and local variables accessed at offsets from ebp
- can avoid using a frame pointer [normally for speed] by accessing parameters and locals variables relative to the stack pointer, but more difficult because the stack pointer can change during execution [BUT easy for a compiler to track]

Positive Parameters

- parameters accessed with +ve offsets from ebp [see stack frame diagram]  
return address at [ebp+4]  
p0 at [ebp+8]  
p1 at [ebp+12]  
...
- local variables accessed with -ve offsets from ebp [see stack frame diagram]

local variable 0 at [ebp-4]

local variable 1 at [ebp-8]

...

# IA32 AND X64

## Consider the IA32 Code for a Simple Function

```
int f (int p0, int p1, int p2) {    // parameters
    int x, y;                      // local variables
    x = p0 + p1;
    ...
    return x + y;                 // result
}
```

- a call `f(p0, p1, p2)` matches stack frame diagram on previous slide
- 3 parameters `p0`, `p1` and `p2`
- 2 local variables `x` and `y`

# IA32 AND X64

## IA32 Code to Call Simple Function...

- parameters  $p0, p1$  and  $p2$  pushed onto stack by caller right to left

$f(1, 2, 3)$

```
push    3          ; push immediate values...
push    2          ; right to left
push    1          ;
call    f          ; call subroutine f
add    esp, 12     ; on return, add 12 to esp to remove parameters
```

push return address  
and jump to f

- inside function parameters accessed relative to ebp [see stack frame diagram]

$p0$  stored at [ebp+8]  
 $p1$  stored at [ebp+12]  
 $p2$  stored at [ebp+16]

# IA32 AND X64

## Accessing Local Variables of Simple Function...

- space allocated on stack for local variables x and y (see diagram)

*x stored at [ebp-4]  
y stored at [ebp-8]*

- $x = p0 + p1$

```
mov  eax, [ebp+8]      ; eax = p0
add  eax, [ebp+12]      ; eax = p0 + p1
mov  [ebp-4], eax       ; x = p0 + p1
```

- return  $x + y$ ;

```
mov  eax, [ebp-4]      ; eax = x
add  eax, [ebp-8]      ; eax = x + y
```

NB: result returned in eax

# IA32 AND X64

## Function Entry

- need instructions on function entry to save ebp [**old frame pointer**]
- initialize ebp [**new frame pointer**]
- allocate space for local variables on stack
- push any non volatile registers used by function onto stack

```
push  ebp          ; save ebp
mov   ebp, esp    ; ebp -> new stack frame
sub   esp, 8       ; allocate space for locals [x and y]
push  ebx        ; save any non volatile registers used by function
<function body>      ; function code
```

NB: cdecl convention means there is NO need to save eax, ecx and edx

# IA32 AND X64

## Function Exit

- need instructions to unwind stack frame at function exit

```
...
pop  ebx          ; restore any saved registers
mov  esp, ebp      ; restore esp  Immediately deallocates local vars
pop  ebp          ; restore previous ebp
ret  0            ; return from function
```

- ret pops return address from stack and...
- adds integer parameter to esp [used to remove parameters from stack]
- if integer parameter not specified, defaults to 0
- NB: cdecl convention - caller removes parameters from stack
- NB: make sure you know why a stack frame needs to be created for each function call  
→ What if you have a recursive function?

→ Need a logical boundary b/w calls  
Environment for each function

# IA32 AND X64

## IA32 Code for Simple Array Accesses

```
int a[100];           // global array of int

main(...) {
    a[1] = a[2] + 3; // constant indices
}
```

NB: int is 4 bytes

NB: a[0] store at address a, a[1] at a+4, a[2] at a+8, a[n] at a+n\*4

```
mov    eax, [a+8];      // eax = a[2]
add    eax, 3            // eax = a[2] + 3
mov    [a+4], eax        // a[1] = a[2] + 3
```

# IA32 AND X64

## IA32 Code for Simple Array Accesses...

```
int p() {  
    int i = ...;                      // local variable i stored at [ebp-4]  
    int j = ...;                      // local variable j stored at [ebp-8]  
    ...  
    a[i] = a[j] + 3;                  // variable indices  
}  
  
mov    eax, [ebp-8]                 // eax = j  
mov    eax, [a+eax*4]               // eax = a[j]  
add    eax, 3                      // eax = a[j]+3  
mov    ecx, [ebp-4]                // ecx = i  
mov    [a+ecx*4], eax             // a[i] = a[j]+3
```

# IA32 AND X64

## Putting it Together - Mixing C/C++ and IA32 Assembly Language

- Example using Visual Studio 2010/2013/2015, VC++ and MASM
- VC++ main(...) calls an assembly language versions of fib(n) to calculate the n<sup>th</sup> Fibonacci number
- create a VC++ Win32 console application
- *right click on project and select "Build Customizations..." and tick masm*
- add fib32.h and fib32.asm files to project [file can be edited within Visual Studio once included, but there doesn't appear to be a way to create an .asm file from within Visual Studio]
- right click on fib32.asm and check [General][Item Type] == Microsoft Macro Assembler
- check project [Properties][Debugger][Debugger Type] == Mixed
- how to look at code generated by VC++ compiler??  
right click on file name [Properties] [C/C++] [Output Files] [Assembler Output] and select Assembly, Machine Code and Source [generates listing file with .cod extension]
- NB: code generated in Debug mode will be different from Release mode

# IA32 AND X64

## Putting it together...

### fib32.h

- declare fib\_IA32a(**int**) and fib\_IA32b(**int**) as external C functions so they can be called from a C/C++ program

```
extern "C" int g;           // external global int
extern "C" int __cdecl fib_IA32a(int); // external function
```

- extern "C" because C++ function names have extra characters which encode their result and parameter types

### fib32.asm

- fib\_IA32a(**int**) – *mechanical* code generation simulating Debug mode
- fib\_IA32b(**int**) – *optimized* code generation simulating Release mode
- NB: MASM specific directives at start of file
- NB: **.data** and **.code** sections
- NB: **public**

# IA32 AND X64

## Mixing C/C++ and IA32 Assembly Language...

IA32codegen.cpp [**main**]

- #include fib32.h
- call fib\_IA32a(n) and fib\_IA32b(n) like any other C/C++ function
- file also contains
  - 1) a C++ version of fib(n) and...
  - 2) a version of fib(n) that mixes C/C++ and IA32 assembly language using the IA32 inline assembler supported by the VC++ compiler
- call ALL versions of fib(n) for n = 1 to 20
- Visual Studio automatically compiles IA32codegen.cpp, assembles fib32.asm and links them to produce an executable which can then be run
- **WARNING:** Visual Studio on SCSS machines (eg ICT Huts) has problems with source files stored on a Network drive

# IA32 AND X64

## x64 Basics

- extension of IA32
- originally developed by AMD
- IA32 registers extended to 64 bits rax ... rsp, rflags and rip
- 8 additional registers r8 .. r15
- 64, 32, 16 and 8 bit arithmetic
- *same* instruction set
- 64 bit virtual and physical address spaces  
[theoretically anyway]
- $2^{64} = 16 \text{ Exabytes} = 16 \times 10^{18} \text{ bytes}$

64			
32			
16			
8			
rax (accumulator)	eax	ax	al
rbx	ebx	bx	bl
rcx	ecx	cl	
rdx	edx	dx	dl
rsi	esi	si	sil
rdi	edi	di	dil
rbp (frame pointer)	ebp	bp	bpl
rsp (stack pointer)	esp	sp	spl
r8	r8d	r8w	r8b
r9	r9d	r9w	r9b
r10	r10d	r10w	r10b
r11	r11d	r11w	r11b
r12	r12d	r12w	r12b
r13	r13d	r13w	r13b
r14	r14d	r14w	r14b
r15	r15d	r15w	r15b

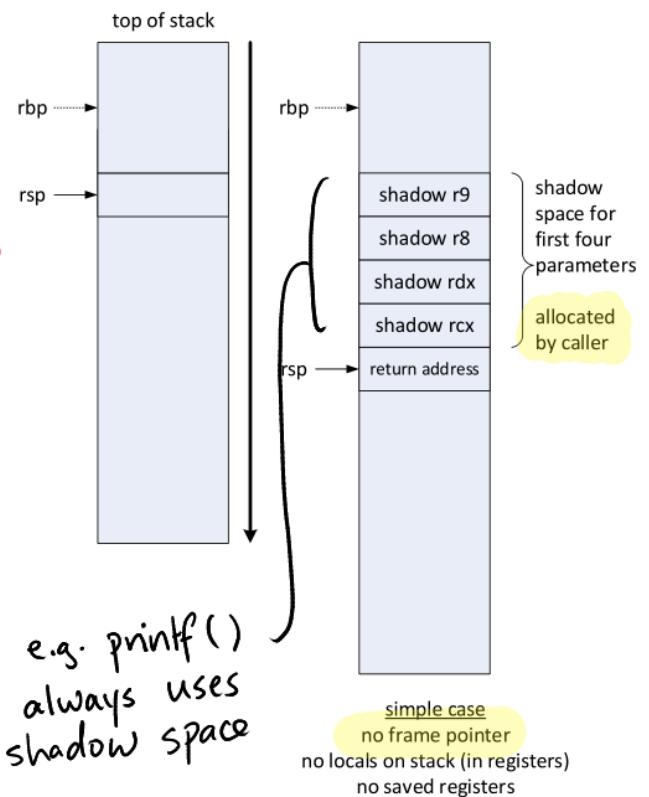
  

rflags [status register]
rip [instruction pointer pc]

# IA32 AND X64

## x64 Function Calling

- use Microsoft calling convention
- first 4 parameters passed in rcx, rdx, r8 and r9 respectively
- additional parameters passed on stack [right to left]
- stack always aligned on an 8 byte boundary
- caller must allocate *shadow space* on stack [conceptually for the parameters passed in rcx, rdx, r9 and r10]
- rax, rcx, rdx, r8, r9, r10 and r11 volatile
- having so many registers often means:
  - can use registers for local variables
  - no need to use a frame pointer
  - no need to save/restore registers



→ Not allocating shadow space risks corrupting the stack

# IA32 AND X64

## x64 Function Calling...

- in the Microsoft x64 calling convention, it is the responsibility of the caller to allocate 32 bytes (4 x 8) of *shadow space* on the stack before calling a function [regardless of the actual number of parameters used] and to deallocate the *shadow space* afterwards
- called functions can use its *shadow space* to spill rcx, rdx, r8, and r9 [spill = save in memory]
- called functions, however, can use its *shadow space* for any purpose whatsoever and consequently may write to and read from it as it sees fit [which is why it needs to be allocated]
- 32 bytes of *shadow space* must be made available to all functions, even those with fewer than four parameters

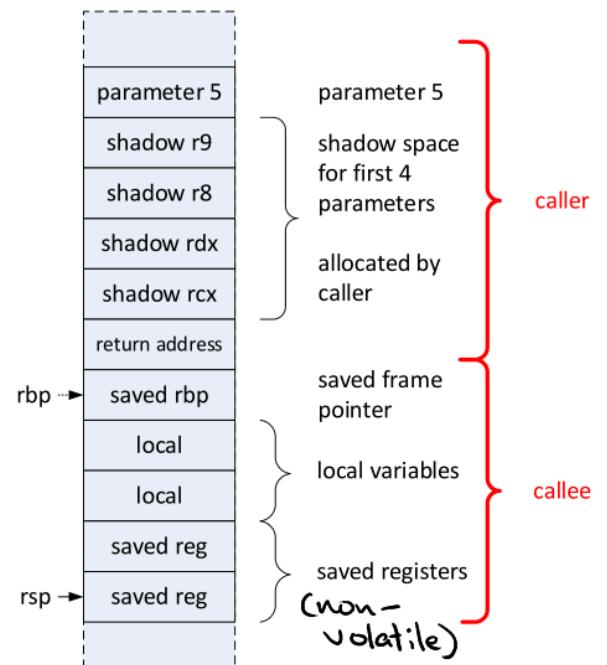
## x64 Function Calling Unix / Linux

- First six parameters are passed  
RDI, RSI, RDX, RCX, R8, R9
- Result returned in RAX

# IA32 AND X64

## x64 Function Calling...

- a more complex x64 stack frame
- callee has 5 parameters, so parameter 5 passed on stack
- parameters 1 to 4 passed in rcx, rdx, r8 and r9
- NB allocated shadow space
- old frame pointer saved and new frame pointer initialised [rbp]
- space allocated for local variables on stack [**if needed**]
- registers saved on stack [**if needed**]



# IA32 AND X64

## x64 Code for a Simple Function

```
_int64 fib(_int64 n) {  
    INT64 fi, fj, t;  
  
    if (n <= 1)  
        return n;  
  
    fi = 0; fj = 1;  
    while (n > 1) {  
        t = fj;  
        fj = fi + fj;  
        fi = t;  
        n--;  
    }  
    return fj;  
}
```

- use \_int64 to declare 64 bit integers [Microsoft specific]
  - alternatively declare 64 bit integers using long long
- #define INT64 long long
- parameter n passed to function in rcx
  - leaf function [as fib doesn't call any other functions]
  - usually easier to code with x64 assembly language rather than IA32 because a simpler stack frame is used and more registers are available

# IA32 AND X64

## x64 Code for a Simple Function...

```
fib_x64:    mov    rax, rcx          ; rax = n
              cmp    rax, 1           ; if (n <= 1)
              jle    fib_x64_1        ; return n
              xor    rdx, rdx          ; fi = 0
              mov    rax, 1           ; fj = 1
fib_x64_0:  cmp    rcx, 1           ; while (n > 1)
              jle    fib_x64_1        ;
              mov    r10, rax          ; t = fj
              add    rax, rdx          ; fj = fi + fj
              mov    rdx, r10          ; fi = t
              dec    rcx              ; n--
              jmp    fib_x64_0        ;
fib_x64_1:  ret                ; return
```

NB: code ONLY uses volatile registers

# IA32 AND X64

## x64 Code for a more Complex Function

```
_int64 xp2(_int64 a, _int64 b) {
    printf("a = %I64d b = %I64d a+b = %I64d\n", a, b, a + b);
    return a + b;      // NB
}
```

- uses %I64d to format a 64 bit integer
- parameters **a** and **b** passed into xp2 in rcx and rdx respectively
- need to call external printf(...) function with 4 parameters

rcx [address of format string]  
rdx [**a**]  
r8 [**b**]  
r9 [**a+b**]

## IA32 AND X64

### x64 Code for a more Complex Function...

```
fpx2 db      'a = %I64d b = %I64d a+b = %I64d',0AH,00H ; ASCII format string
                                         |          |
                                         newline    null terminator
xp2: push   rbx           ; save rbx
       sub    rsp, 32        ; allocate shadow space
• lea    r9, [rcx+rdx]    ; printf parameter 4 in r9 {a+b}
       mov    r8, rdx        ; printf parameter 3 in r8 {b}
       mov    rdx, rcx        ; printf parameter 2 in rdx {a}
• lea    rcx, fpx2       ; printf parameter 1 in rcx {&fpx2}
       mov    rbx, r9         ; save r9 in rbx so preserved across call to printf
       call   printf          ; call printf
       mov    rax, rbx        ; function result in rax = rbx {a+b}
       add    rsp, 32         ; deallocate shadow space
       pop    rbx            ; restore rbx
       ret
```

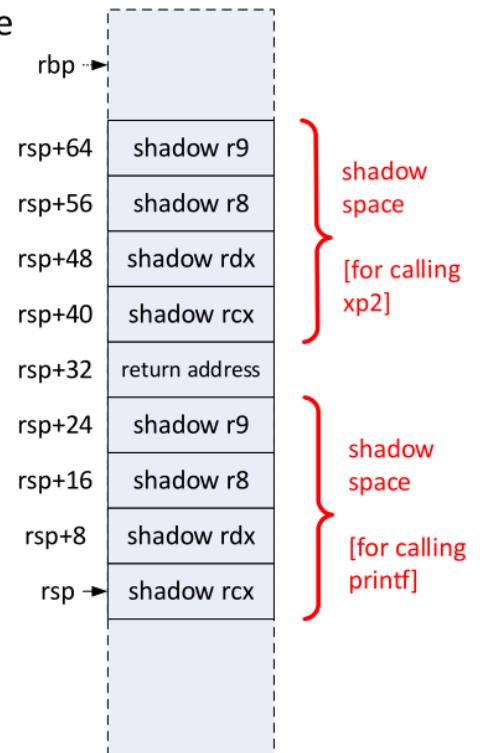
- LEA used for storing memory addresses (pointers)
  - ⇒ De-referencing values [rcx+rdx] is the same as a mov
    - ↳ Adds ability to add multiple operands at once
- Put pointer to address of string in rcx

# IA32 AND X64

## x64 Code for a more Complex Function...

- instead of using rbx to preserve r9 across the call to printf,  
an alternate approach is to use a location its shadow space  
**[eg. rsp+64]**

```
xp2: sub    rsp, 32      ; allocate shadow space
      lea     r9, [rcx+rdx] ; printf parameter 4 in r9 {a+b}
      mov     r8,rdx       ; printf parameter 3 in r8 {b}
      mov     rdx,rcx       ; printf parameter 2 in rdx {a}
      lea     rcx,fxp2     ; printf parameter 1 in rcx
      mov     [rsp+64],r9   ; save r9 in shadow space so...
      call    printf        ; preserved across call to printf
      mov     rax,[rsp+64]  ; result in rax = saved r9 {a+b}
      add    rsp,32        ; deallocate shadow space
      ret                 ; return
```



# IA32 AND X64

## x64 Code for a more Complex Function...

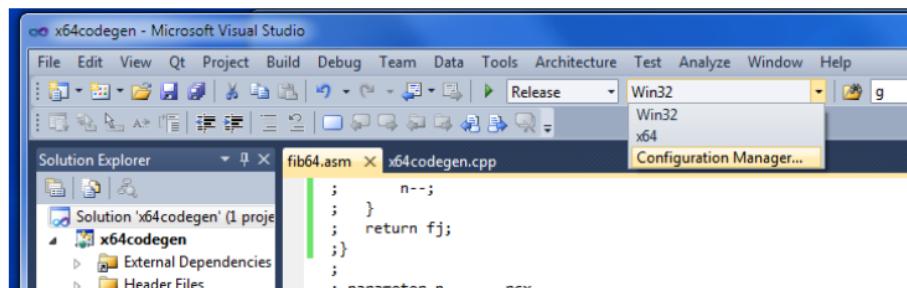
Typical code generation strategy

- shadow space allocated ONCE at start of function
- allocate enough stack space to accommodate calls to the function with the most parameters [NB: must allocate a minimum 32 bytes for the shadow space]
- use the same stack space [and registers] to pass parameters to ALL the functions it calls
- straightforward for compiler to determine how much stack space is required

# IA32 AND X64

## x64 Code for a Simple Function...

- fib64.h, fib64.asm and x64codegen.cpp on CS3021/3421 website
- need to create a console application and use the Configuration Manager to select a x64 solution platform



- fib64.asm
  - extern printf:near (for printf) ; allows external printf() function to be called
- ALL OK with VS2013, but printf linker problems with VS2015 [~~solve temporarily by configuring VS2015 to use VS2013 toolset for project~~]
- no x64 inline assembler, can use intrinsics defined in [instrin.h](#) instead

# IA32 AND X64

## Summary

- you are now able to:
  - write simple IA32 assembly language functions
  - write simple x64 assembly language functions
  - call IA32/x64 assembly language functions from C/C++
  - program the two most widely used CPUs in assembly language  
**[IA32/x64 and ARM]**



## Study Points

- Eax = Accumulator - register for storing intermediate arithmetic ops // store return value
- Ebx = Base Index - for use with arrays // non-volatile
- ECX = Counter - Loops and strings // store 1st param

• Edx = Extend precision of Eax // storing short term vars //  
store 2nd param

• Esi = Source pointer

• Edi = Destination pointer

• Function calling conventions for x86 and x64  
→ parameter passing

→ caller / callee jobs

→ volatile / non-volatile registers

→ Shadow space

• Stack frame diagrams